

Summary and Highlights

Congratulations! You have completed this module. At this point in the course, you know:

- Instruction-tuning involves training models with expert-curated datasets.
- For instruction-tuning, the model requires instructions and answers.
- Instruction-tuning helps perform a wide variety of tasks by interpreting and executing instructions more effectively.
- Instruction-tuning uses three components: Instructions, input, and output.
- Adjust the prompt format to maintain compatibility with the different models' tokenizers.
- Instruction masking focuses on the loss calculation of specific tokens.
- Load a dataset using CodeAlpaca 20k dataset.
- Format the dataset using the *formatting_prompts_func* function.
- For creating a formatted dataset, use two code blocks for generating instructions with and without responses.
- To create a model, fine-tune Facebook's opt-350m model.
- Define the collator using *DataCollatorForCompletionOnlyLM* to prepare data batches for training language models.
- Define the trainer by creating the *SFTTrainer* object.
- Generate a text pipeline from the Transformers library.
- Evaluate the model's text generation using the BLEU score.
- The reward model takes prompt as an input and response as an output regarding reward or score.
- Reward modeling helps in quantifying quality responses, guiding model optimization, incorporating reward preferences, and ensuring consistency and reliability of the responses.
- The scoring function takes the query and appends the chatbot's responses.
- The dataset synthetic-instruct-gptj-pairwise from Hugging Face trains and evaluates instruction-following models.

- Defining the *preprocess_function()* helps format the keys and tokenize the data for the reward trainer.
- *TrainingArguments* class from the Transformers library defines the training arguments.
- Reward trainers orchestrate the process and save and evaluate the model using the *trainer.train()* method.
- The tokenizing process generates scores and compares the output of two functions to achieve the desired win rate.
- Reward model training is an advanced technique that trains a model to identify desired outputs generated by another model and assign scores to the outcome based on its relevance and accuracy.
- Training the scoring function helps generate rewards effectively.
- Generating reward model loss, the encoder model generates responses as contextual embeddings.
- Using the Bradley–Terry reward loss model, you can understand the reward loss model by generating the cost or loss function.